

# Analyzing Genre Influence on Spotify Track Characteristics

Necessary R libraries:

```
library(base)
library(reticulate) #for Python interface
library(psych) #for skew/kurt functions
library(moments)
library(stats)
library(lsr)
```

Necessary Python libraries:

```
import kaggle
import requests
import kaggle.cli
import sys
import numpy
import pandas as pd
from pathlib import Path
from zipfile import ZipFile
from kaggle.cli import main
```

```
#download data set from Kaggle
dataset="thedeastator/spotify-tracks-genre-dataset"
sys.argv=[sys.argv[0]] + f"datasets download {dataset}".split(" ")

kaggle_zip=ZipFile(f"{dataset.split('/')[1]}.zip")
songdat={i.filename:pd.read_csv(kaggle_zip.open(i)) for i in kaggle_zip.infolist()}["train.csv"]

#Close the ZipFile object
kaggle_zip.close()
```

Clean data (python): group subgenres into broader categories for new genre row

```

import numpy

#Pop
songdat.loc[songdat['track_subgenre'].str.contains('pop', case=False, na=False) | songdat
['track_subgenre'].str.contains('idol', case=False, na=False), 'track_genre']='Pop'

#Metal
songdat.loc[songdat['track_subgenre'].str.contains('metal', case=False, na=False) | (son
gdat['track_subgenre']=='hardcore'), 'track_genre']='Metal'

#Rock
rock_conditions=songdat['track_subgenre'].str.contains('rock', case=False, na=False) | s
ongdat['track_subgenre'].isin(['goth', 'emo', 'garage', 'grunge', 'punk'])
songdat.loc[rock_conditions, 'track_genre']='Rock'

#Reggae
songdat.loc[songdat['track_subgenre'].str.contains('reggae', case=False, na=False) | (so
ngdat['track_subgenre']=='ska'), 'track_genre']='Reggae'

#Electronic
electronic_conditions=songdat['track_subgenre'].str.contains('dance', case=False, na=Fa
lse) | songdat['track_subgenre'].str.startswith('electro') | songdat['track_subgenre'].st
r.startswith('dub') | songdat['track_subgenre'].str.contains('house', case=False, na=Fa
lse) | songdat['track_subgenre'].str.contains('techno', case=False, na=False) | songdat
['track_subgenre'].isin(['edm', 'party', 'dubstep', 'hardstyle', 'breakbeat', 'club', 'd
rum-and-bass', 'idm', 'trance', 'trip-hop', 'industrial', 'happy'])
songdat.loc[electronic_conditions, 'track_genre']='Electronic'

#Disco
songdat.loc[songdat['track_subgenre'].str.contains('disco', case=False, na=False), 'trac
k_genre']='Disco'

#Jazz
songdat.loc[songdat['track_subgenre'].str.contains('jazz', case=False, na=False), 'track
_genre']='Jazz'

#Blues
songdat.loc[songdat['track_subgenre']=='blues', 'track_genre']='Blues'

#Folk
songdat.loc[songdat['track_subgenre'].isin(['folk', 'bluegrass']), 'track_genre']='Folk'

#Country
songdat.loc[songdat['track_subgenre']=='country', 'track_genre']='Country'

#R&B/Soul
songdat.loc[songdat['track_subgenre'].isin(['r-n-b', 'funk', 'soul', 'groove', 'afrobea
t']), 'track_genre']='R&B/Soul'

#Latin
latin_conditions=songdat['track_subgenre'].str.contains('latin', case=False, na=False) |
songdat['track_subgenre'].isin(['brazil', 'forro', 'mpb', 'pagode', 'salsa', 'samba', 's

```

```

ertanejo', 'spanish'])
songdat.loc[latin_conditions, 'track_genre']='Latin'

#World
world_conditions=songdat['track_subgenre'].str.contains('world', case=False, na=False) |
songdat['track_subgenre'].isin(['french', 'german', 'indian', 'malay', 'swedish', 'turkish', 'british', 'anime'])
songdat.loc[world_conditions, 'track_genre']='World'

#Hip-hop
songdat.loc[songdat['track_subgenre'].isin(['hip-hop', 'sad']), 'track_genre']='Hip-hop'

#Kids
songdat.loc[songdat['track_subgenre'].isin(['children', 'kids', 'disney', 'comedy', 'show-tunes']), 'track_genre']='Kids/Family'

#Alternative
songdat.loc[songdat['track_subgenre'].isin(['alternative', 'chill', 'indie', 'new-age', 'sleep', 'ambient']), 'track_genre']='Alternative'

#Instrumental
songdat.loc[songdat['track_subgenre'].isin(['acoustic', 'guitar', 'piano']), 'track_genre']='Instrumental'

#Singer-songwriter
songdat.loc[songdat['track_subgenre'].str.contains('songwriter', case=False, na=False), 'track_genre']='Singer-songwriter'

#Classical
songdat.loc[songdat['track_subgenre'].isin(['classical', 'gospel', 'opera']), 'track_genre']='Classical'

#Capitalize first letter in Track Subgenre
songdat['track_subgenre']=songdat['track_subgenre'].apply(lambda x: x.capitalize())

```

Check if we've left behind any track\_subgenre rows that haven't been matched with track\_genre

```

if(songdat['track_genre'].isnull().sum())!=0:
    null_rows=songdat[songdat['track_genre'].isnull()]
    grouped_null=null_rows.groupby('track_subgenre').size().reset_index(name='count')
    print(grouped_null)

```

```

##   track_subgenre  count
## 0      Grindcore   1000
## 1    Honky-tonk   1000
## 2      Iranian   1000
## 3      Romance   1000
## 4        Study   1000
## 5         Tango   1000

```

```

#Missed Grindcore, Honky-tonk, Iranian, Romance, Study, Tango

#Add Grindcore to Metal
songdat.loc[songdat['track_subgenre']=='Grindcore', 'track_genre']='Metal'

#Add Honky-tonk to Country
songdat.loc[songdat['track_subgenre']=='Honky-tonk', 'track_genre']='Country'

#Add Iranian to World
songdat.loc[songdat['track_subgenre']=='Iranian', 'track_genre']='World'

#Add Tango to Latin
songdat.loc[songdat['track_subgenre']=='Tango', 'track_genre']='Latin'

#Add Study to Alternative
songdat.loc[songdat['track_subgenre']=='Study', 'track_genre']='Alternative'

#Not sure where Romance should go, let's look at the artist names
romance=songdat[songdat['track_subgenre']=='Romance']
print(romance[['artists', 'track_name']].sort_values(by='artists'))

```

```

##                artists                track_name
## 93700      Alexander Vertinsky      Lilovyy negr
## 93279      Alexander Vertinsky  Ja segodnja smejus' nad soboy
## 93312      Alexander Vertinsky  Ja segodnja smejus' nad soboy
## 93416      Alexander Vertinsky      Pesenka o zhene
## 93457      Alexander Vertinsky      Pesenka o zhene
## ...                ...                ...
## 93802      Тамара Церетели      Взгляд твоих чёрных очей
## 93405  Татьяна Комова;Георгий Квик      Солнышко
## 93859  Татьяна Комова;Георгий Квик      Пятеро детей
## 93454  Татьяна Комова;Георгий Квик      Солнышко
## 93553  Татьяна Комова;Георгий Квик      Сумерки
##
## [1000 rows x 2 columns]

```

```

#Looks like Romance is a form of classical Russian music, add to World
songdat.loc[songdat['track_subgenre']=='Romance', 'track_genre']='World'

```

Choose numeric variables we want to compare against track\_genre, our predictor variable. Let's try "popularity", "energy", "tempo", "duration\_ms", "danceability", "loudness". We'll scale all our raw values first to make comparison easier.

```

# Convert py$songdat from Python to R dataframe

numeric_columns <- c("popularity", "energy", "tempo", "duration_ms", "danceability","loudness")

#Quick function to scale our rows quickly
scale_column <- function(df, column_name) {
  if (!column_name %in% colnames(df)) {
    stop("Column not found in the songdat")
  }

  df_scaled <- df
  df_scaled[[paste0(column_name, "_scaled")]] <- scale(df[[column_name]], scale=FALSE)

  return(df_scaled)
}

#cycle through all our rows in numeric_columns, scaling each, and append to songdat
for (column_name in as.character(numeric_columns)) {
  scaled_column<-scale(py$songdat[column_name],scale=FALSE)
  py$songdat[[paste0(column_name, "_scaled")]] <-scaled_column[, 1]
}

```

Now check the Skew/Kurtosis values for all our dv's of interest.

```

#Look at skew/kurt for scaled variables
skewness_scaled <- skew(py$songdat[paste0(numeric_columns,"_scaled")])
kurtosis_scaled <- kurtosis(py$songdat[paste0(numeric_columns,"_scaled")])

#Combine results into a data frame for easier viewing
skew_kurtosis_df <- data.frame(skewness=skewness_scaled, kurtosis=kurtosis_scaled)
print(skew_kurtosis_df)

```

##	skewness	kurtosis
## popularity_scaled	0.04640129	2.072233
## energy_scaled	-0.59698571	2.474260
## tempo_scaled	0.23228875	2.891372
## duration_ms_scaled	11.19488687	357.936795
## danceability_scaled	-0.39948612	2.815453
## loudness_scaled	-2.00648913	8.895967

```

#Duration_ms has way too much skew/kurt, 11.19503417 and 357.936795! Not worth attempting
a transformation, too much variance
numeric_columns<-subset(numeric_columns, numeric_columns!="duration_ms") #Drop Duration_ms

```

All variables look good except for 2, duration\_ms and loudness. We removed duration\_ms because the skew/kurt was to extreme to fix through transformations. Loudness has a kurtosis of 8.9, which is problematic but might benefit from transformations. Let's try square root, inverse, and log transformations.

```
song_pos=abs(py$songdat$loudness+1e-10)#There are 0s, to avoid NaN results add a really small number so it's technically positive
```

```
#Square Root
```

```
songNormSqt <- data.frame((song_pos+1)^0.5)
```

```
skewnessSqt <- skew(song_pos)
```

```
kurtosisSqt <- kurtosis(song_pos)
```

```
skewnessSqt
```

```
## [1] 2.009743
```

```
kurtosisSqt
```

```
## [1] 8.902822
```

```
#Inverse
```

```
songNormIn <- 1/(song_pos+1)
```

```
skewnessIn <- skew(songNormIn)
```

```
kurtosisIn <- kurtosis(songNormIn)
```

```
skewnessIn
```

```
## [1] 2.606847
```

```
kurtosisIn
```

```
## [1] 19.33271
```

```
#Log
```

```
songNormLg <- log10(song_pos+1)
```

```
skewnessLg <- skew(songNormLg)
```

```
kurtosisLg <- kurtosis(songNormLg)
```

```
kurtosisLg #Log looks good, kurtosis improved from 8.9 to 3.6 and skew is still minimal at .20.
```

```
## [1] 3.586076
```

```
skewnessLg
```

```
## [1] 0.1690073
```

The Log transformation for loudness looks great, skew dropped to 3.6 from 8.9. Let's apply this normally distributed version and use it as our loudness variable.

```
#Let's replace our raw loudness data with our more normalized log transformation
#use unlist to ensure songNormLg is a 1-dim array
loudness_log_scaled_standardized <- scale(unlist(songNormLg)) #rescale

#assign scaled loudness values to songdat
py$songdat$loudness_log_scaled <- loudness_log_scaled_standardized[,1]

#update the numeric column list since we aren't using the raw
#Find the index of "loudness"
index_to_replace <- which(numeric_columns=="loudness")

#replace "loudness" with "loudness_log_scaled"
numeric_columns[index_to_replace] <- "loudness_log_scaled"
```

Let's perform a series of one-way ANOVAs, using our numeric\_column ("popularity", "energy", "tempo", "danceability", "loudness\_log\_squared") as our dv's, and track\_genre as our iv

```
for(column_name in as.character(numeric_columns)){
  print(column_name)
  aovTemp<-aov(py$songdat[[column_name]]~py$songdat$track_genre)
  print(summary(aovTemp)) #Print ANOVA results
  print(etaSquared(aovTemp,type=3,anova=TRUE)) #...and eta-squared (effect size)
}
```

```
## [1] "popularity"
##
##           Df    Sum Sq Mean Sq F value Pr(>F)
## py$songdat$track_genre    18  3043694  169094    359.1 <2e-16 ***
## Residuals    113981  53672693    471
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
##           eta.sq eta.sq.part      SS      df      MS
## py$songdat$track_genre 0.05366515 0.05366515  3043694    18 169094.0867
## Residuals    0.94633485          NA  53672693 113981   470.8916
##
##           F      p
## py$songdat$track_genre 359.0935 0
## Residuals          NA NA
## [1] "energy"
##
##           Df    Sum Sq Mean Sq F value Pr(>F)
## py$songdat$track_genre    18    2146   119.22    2682 <2e-16 ***
## Residuals    113981    5066    0.04
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
##           eta.sq eta.sq.part      SS      df      MS
## py$songdat$track_genre 0.2975362 0.2975362  2145.938    18 119.2187990
## Residuals    0.7024638          NA  5066.422 113981   0.0444497
##
##           F      p
## py$songdat$track_genre 2682.106 0
## Residuals          NA NA
## [1] "tempo"
##
##           Df    Sum Sq Mean Sq F value Pr(>F)
## py$songdat$track_genre    18  3158697  175483    201.4 <2e-16 ***
## Residuals    113981  99291326    871
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
##           eta.sq eta.sq.part      SS      df      MS
## py$songdat$track_genre 0.03083159 0.03083159  3158697    18 175483.1472
## Residuals    0.96916841          NA  99291326 113981   871.1217
##
##           F      p
## py$songdat$track_genre 201.445 0
## Residuals          NA NA
## [1] "danceability"
##
##           Df    Sum Sq Mean Sq F value Pr(>F)
## py$songdat$track_genre    18    680   37.78    1564 <2e-16 ***
## Residuals    113981    2753    0.02
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
##           eta.sq eta.sq.part      SS      df      MS
## py$songdat$track_genre 0.1980524 0.1980524  679.9724    18 37.77624624
## Residuals    0.8019476          NA  2753.3225 113981   0.02415598
##
##           F      p
## py$songdat$track_genre 1563.847 0
## Residuals          NA NA
## [1] "loudness_log_scaled"
##
##           Df    Sum Sq Mean Sq F value Pr(>F)
## py$songdat$track_genre    18  22528  1251.5    1560 <2e-16 ***
## Residuals    113981  91471    0.8
```



```
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##               eta.sq eta.sq.part      SS      df      MS
## py$songdat$track_genre 0.1976133  0.1976133 22527.72    18 1251.5398287
## Residuals           0.8023867           NA 91471.28 113981    0.8025134
##               F      p
## py$songdat$track_genre 1559.525  0
## Residuals              NA NA
```

Surprisingly all of our dv's produced significant results, meaning track popularity, energy, tempo, danceability, and loudness are all influenced by a song's genre,  $p < .001$ . The largest effect sizes were seen for danceability ( $\eta^2 = .24$ ) and loudness ( $\eta^2 = .22$ ).

Since we know the means differ between track genre's, lets do some summary stats for each of our iv's by track genre. Pairwise analysis would be too much here since we have 18 distinct genre values, but we know we can trust our mean/median/modes due to significant ANOVA results.

```
options(max.print = 1000000)

#print the mean, median, sd for all our values of interest
summary_stats_by_genre <- aggregate(py$songdat[numeric_columns], by=list(track_genre=py
$songdat$track_genre), FUN=function(x) c(mean=mean(x), median=median(x), sd=sd(x)))

print(summary_stats_by_genre)
```

##	track_genre	popularity.mean	popularity.median	popularity.sd		
## 1	Alternative	35.90257	39.00000	23.80630		
## 2	Blues	31.18800	34.00000	27.48425		
## 3	Classical	26.43833	25.00000	17.95328		
## 4	Country	16.69150	12.00000	20.29929		
## 5	Disco	33.52200	32.00000	24.75695		
## 6	Electronic	30.84392	27.00000	21.80143		
## 7	Folk	31.84100	27.00000	18.94955		
## 8	Hip-hop	45.06900	55.00000	25.47522		
## 9	Instrumental	39.09400	40.50000	17.88250		
## 10	Jazz	13.62800	0.00000	23.18290		
## 11	Kids/Family	26.66900	23.00000	14.65071		
## 12	Latin	34.37182	40.00000	18.85740		
## 13	Metal	31.25086	25.00000	19.44254		
## 14	Pop	41.35000	45.00000	22.64844		
## 15	R&B/Soul	30.20560	32.00000	23.36409		
## 16	Reggae	26.53800	33.00000	26.25452		
## 17	Rock	37.63000	37.00000	22.84109		
## 18	Singer-songwriter	37.81300	43.00000	27.68428		
## 19	World	33.22636	38.00000	21.92392		
##	energy.mean	energy.median	energy.sd	tempo.mean	tempo.median	tempo.sd
## 1	0.4153339	0.3960000	0.2822690	110.32292	107.99200	35.60832
## 2	0.5818775	0.5835000	0.2205792	116.56835	114.54000	30.70551
## 3	0.3610459	0.3270000	0.2490275	113.05316	113.26450	30.72567
## 4	0.4818810	0.4530000	0.2230807	120.26009	117.99700	30.84552
## 5	0.7375650	0.7770000	0.1877227	121.97438	123.98100	19.22186
## 6	0.7555168	0.7880000	0.1769618	127.68878	125.11850	25.61376
## 7	0.5380433	0.5445000	0.2130263	122.52417	120.96550	27.67009
## 8	0.5725001	0.5825000	0.1981201	117.91727	112.46150	29.98537
## 9	0.3601566	0.3030000	0.2630827	116.36518	114.78550	31.72331
## 10	0.3529544	0.3320000	0.1858928	112.63647	109.15200	31.64017
## 11	0.5017937	0.5010000	0.2549701	114.50411	112.95550	30.66556
## 12	0.6681083	0.7010000	0.1892528	122.81746	119.03650	30.74957
## 13	0.8859689	0.9370000	0.1318622	126.01259	123.00800	29.89331
## 14	0.6471027	0.6610000	0.2167343	124.35875	123.88100	29.40967
## 15	0.6542870	0.6750000	0.1985774	119.87530	119.99300	26.63007
## 16	0.7525510	0.7680000	0.1347853	122.35945	109.98400	32.93197
## 17	0.7165144	0.7620000	0.2094055	126.45852	125.00300	30.45898
## 18	0.4341884	0.4360000	0.2064300	119.73732	117.98200	30.94120
## 19	0.5507639	0.5640000	0.2573138	118.22343	118.32300	29.92687
##	danceability.mean	danceability.median	danceability.sd			
## 1	0.4830041	0.5150000	0.2253829			
## 2	0.5685670	0.5790000	0.1472112			
## 3	0.3895946	0.3890000	0.1494489			
## 4	0.5632585	0.5640000	0.1203286			
## 5	0.6766920	0.6930000	0.1232074			
## 6	0.6349987	0.6470000	0.1416097			
## 7	0.5466030	0.5520000	0.1195352			
## 8	0.7142660	0.7270000	0.1234201			
## 9	0.5149096	0.5260000	0.1494902			
## 10	0.5099750	0.4990000	0.1413249			
## 11	0.5998361	0.6070000	0.1948149			

## 12	0.6200453	0.6300000	0.1386085
## 13	0.3999535	0.4000000	0.1538802
## 14	0.5768700	0.5810000	0.1365507
## 15	0.6292348	0.6470000	0.1559566
## 16	0.6948427	0.7210000	0.1370737
## 17	0.5161622	0.5200000	0.1410385
## 18	0.5620220	0.5660000	0.1298911
## 19	0.5263003	0.5400000	0.1893377
##	loudness_log_scaled.mean	loudness_log_scaled.median	loudness_log_scaled.sd
## 1	0.9242381997	0.8860563633	1.1442901784
## 2	0.1644425561	0.1563732667	0.7827120462
## 3	0.8985326803	0.8000491275	1.1389713908
## 4	0.3778392010	0.4944670859	0.7724719383
## 5	-0.1260990675	-0.0845407694	0.8825831523
## 6	-0.3260947153	-0.3264271720	0.9287134098
## 7	0.3923454423	0.4004746721	0.6957829603
## 8	0.0451322918	-0.0002549701	0.8444869827
## 9	0.8976135603	0.8737704039	1.1029524403
## 10	0.7700157736	0.7353991642	0.6959407367
## 11	0.5356254782	0.5076108818	0.9239830595
## 12	-0.1986341967	-0.1961881537	0.7740651700
## 13	-0.5797516891	-0.6225886092	0.7550342060
## 14	-0.1764631239	-0.1049569029	0.8789564985
## 15	-0.0994416018	-0.0750636006	0.7329113067
## 16	-0.5869576556	-0.6137241482	0.6817710769
## 17	-0.2283807470	-0.2486230457	0.8347214506
## 18	0.4444801848	0.4755011094	0.7132964409
## 19	0.3202800714	0.2667884074	1.0009273097

Looks good, let's export our summary data into a CSV to import into Tableau!

```
write.csv(summary_stats_by_genre, file="summary_stats_by_genre.csv", row.names=FALSE)
```